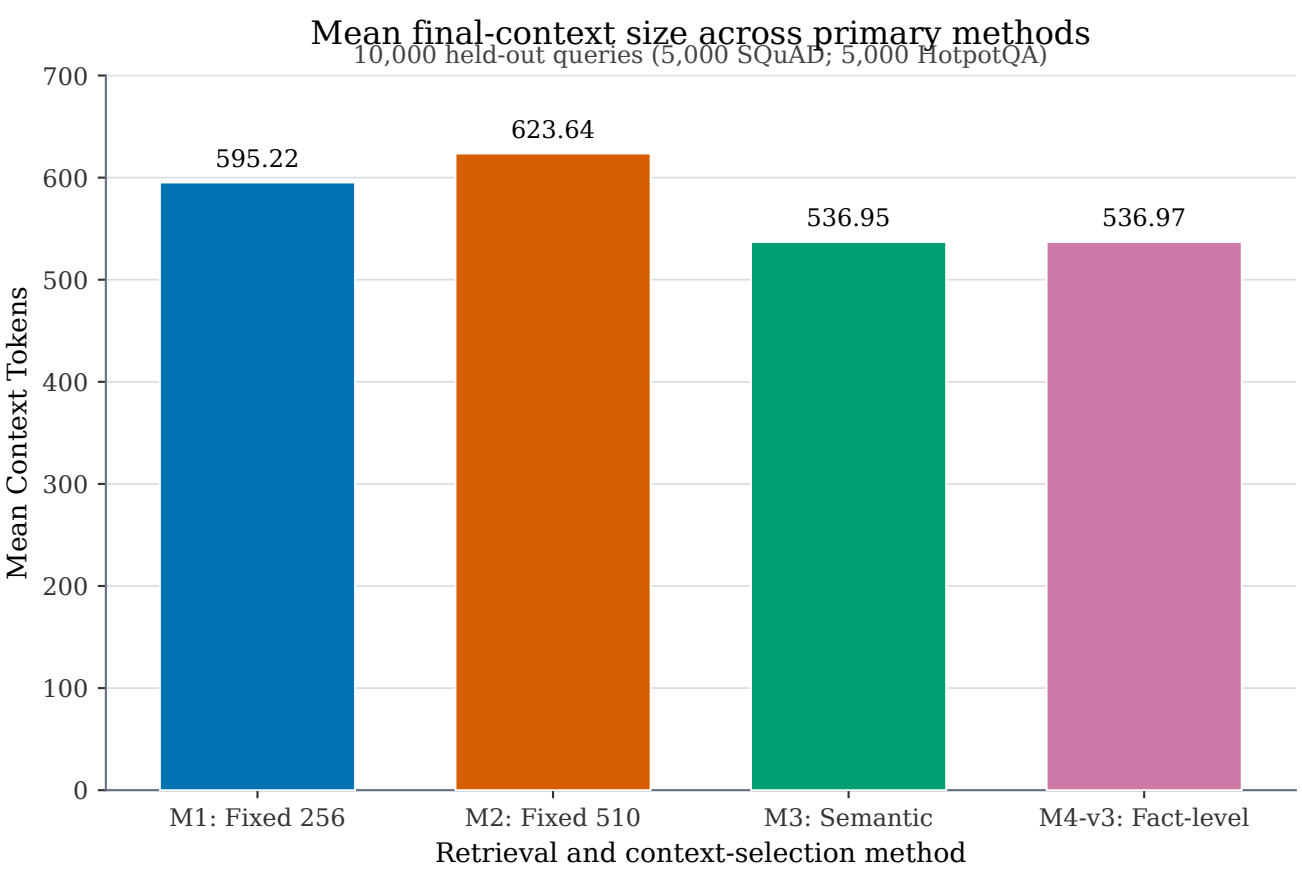
Mean final-context size across primary methods
10,000 held-out queries (5,000 SQuAD; 5,000 HotpotQA)
Mean Context Tokens
700
600
500
400
300
200
100
0
595.22
623.64
536.95
536.97
M1: Fixed 256
M2: Fixed 510
M3: Semantic
M4-v3: Fact-level
Retrieval and context-selection method
A smaller context does not automatically imply better evidence coverage; interpret alongside Recall@k.